

33.66 GB ONE-TIME DOWNLOADS	1.68 GB AVG PER STUDENT	385.78 MB ASSISTANT, WHOLE ROOM	0.143 Mbps ASSISTANT ROOM AVERAGE	0.39 Mbps ALL STREAMS AT ONCE
--	-----------------------------------	--	--	--

VENUE IMPACT

AI assistant traffic is negligible: 385.78 MB for the whole room over the full session, averaging 0.143 Mbps — less than one video call. Plan around the one-time 33.66 GB download burst: stagger the install step or ask attendees to install before arriving.

INSTALLERS — 31.28 GB TOTAL

ITEM	SIZE	STUDENTS	AGGREGATE	
Python (macOS) python-3.14.7-macos11.pkg	78.56 MB	10	785.58 MB	
Python (Windows) python-3.14.7-amd64.exe	33.26 MB	10	332.58 MB	
R (macOS) R-4.6.1-arm64.pkg	105.07 MB	10	1.05 GB	
R (Windows) R-4.6.1-win.exe	91.75 MB	10	917.50 MB	
Positron (macOS) Positron-2026.08.1-2-arm64.dmg	1.05 GB	10	10.45 GB	
Positron (Windows) Positron-2026.08.1-2-UserSetup-x64.exe	522.83 MB	10	5.23 GB	
RStudio Desktop (macOS) RStudio-2026.08.0-187.dmg	852.53 MB	10	8.53 GB	
RStudio Desktop (Windows) RStudio-2026.08.0-187.exe	399.46 MB	10	3.99 GB	

WORKSHOP PACKAGES — 2.37 GB TOTAL



ITEM	SIZE	STUDENTS	AGGREGATE	
Python packages: querychat + shinychat 71 packages resolved from querychat, shinychat	67.81 MB	20	1.36 GB	
R packages: querychat + shinychat 48 packages resolved from querychat, shinychat	50.73 MB	20	1.01 GB	

AI ASSISTANT TRAFFIC

ACTIVITY	CONVS/STUDENT	PER CONVERSATION	PER STUDENT	ROOM TOTAL
First conversation with Posit Assistant	1	2.68 MB	2.68 MB	53.57 MB
querychat exercises (natural-language data queries)	2	321.84 KB	643.68 KB	12.87 MB
General assistant use (debugging, plots, dashboards)	4	3.99 MB	15.97 MB	319.33 MB

Whole room, whole session: 385.78 MB — an average of 0.143 Mbps. If every student’s assistant streamed a response at the same instant: about 0.39 Mbps, less than one video call.

WORST CASE — EVERY STUDENT DOWNLOADS EVERYTHING AT THE VENUE, AT ONCE

VENUE LINK	ALL DOWNLOADS COMPLETE IN
50 Mbps	128.2 min 
100 Mbps	64.1 min 
250 Mbps	25.6 min 
500 Mbps	12.8 min 
1000 Mbps	6.4 min 

Assumes shared Wi-Fi delivers about 70% of the advertised link speed to actual downloads. Staggering the install step or asking attendees to install before arriving removes this burst entirely.

ASSUMPTIONS

Download sizes are measured live from the release CDNs and package indexes (largest platform wheel per package, conservative). Assistant traffic is modeled: 4 bytes per request token, 30 bytes per streamed token including SSE framing, 1.08× TLS/HTTP overhead, responses streaming at 75 tokens/second. Chat APIs are stateless — each round trip re-sends the conversation so far, which the model accounts for; provider-side prompt caching reduces cost and latency but not network transfer. Full model description and grounding: github.com/statik/measuring-agent-bandwidth (METHODOLOGY.md).